

A CONSTRUCTION OF THE EXTREMAL TORUS FIXED POINT, AND A CRITERION FOR SIMPLE SPECTRUM

ABSTRACT. Let Γ be a simply-laced Dynkin diagram, $\lambda = \omega_\kappa$ a fundamental weight and w an element of the Weyl group. When $\mu = w\lambda$ the Nakajima quiver variety $\mathcal{M}(v, w)$ with $\lambda - \mu = \sum_i v_i \alpha_i$ is a single point, carrying a unique T -fixed representation $V = \bigoplus_i V_i$. Dinkins, Karpov and Krylov raise, in §5.4.2 of [4], the problem of characterising combinatorially those w for which T acts on every V_i with simple spectrum, in which case the fixed point is encoded by a coloured poset and the vertex function is a sum over reverse plane partitions.

We solve this by producing V itself. Let Π be the preprojective algebra of Γ and I_κ the injective envelope of the simple S_κ . Reading a reduced word for w from the right and recording at each step the integer $c = \langle \alpha_i^\vee, \mu_{\text{cur}} \rangle$ produces a sequence σ of vertices with multiplicities. We prove that

$$V \cong \text{soc}_\sigma(I_\kappa),$$

the iterated socle of Geiss–Leclerc–Schröer taken along σ . The proof is uniform in type and rests on an identity, derived from Crawley-Boevey’s formula for Ext^1 over a preprojective algebra, together with a squeeze against $\dim I_\kappa = \lambda - w_0 \lambda$. Simple spectrum then becomes the assertion that no two elements of the resulting coloured graded set share a colour and a degree, which is decidable in polynomial time from any reduced word. In type D_n , for non-spin λ , we make the answer completely explicit: writing $\mu = \sum_{i \in S} \eta_i e_i$ in Bourbaki coordinates, the spectrum is simple if and only if $r + 2m \leq 2$, where m counts the minus signs among the first $n - 2$ coordinates and $r = |S \cap \{n - 1, n\}|$. This last statement rests, for $\lambda = \omega_k$ with $k \geq 2$, on a truncation property of I_κ that we verify by machine rather than prove; the main theorem above, and the reduction of the criterion to the trivalent node, do not. We also record an element of D_4 that is fully commutative but does not have simple spectrum, showing that the sufficient condition considered in [4, Ex. 5.12] fails in both directions.

1. INTRODUCTION

1.1. The problem. Fix a simply-laced Dynkin diagram Γ with vertex set I , Cartan matrix $C = (a_{ij})$, simple roots α_i , fundamental weights ω_i and Weyl group W . Let $\mathcal{M}(v, w)$ denote the Nakajima quiver variety attached to Γ with dimension vector v and framing w , and write $\lambda = \sum_i w_i \omega_i$, $\mu = \lambda - \sum_i v_i \alpha_i$. When μ is an extremal weight of the irreducible representation of highest weight λ — that is, $\mu = w\lambda$ for some $w \in W$ — the variety $\mathcal{M}(v, w)$ is a single point [7, 8], and that point is a representation $V = \bigoplus_{i \in I} V_i$ of the doubled quiver satisfying the moment map equations, unique up to isomorphism.

The point is fixed by the torus T acting on the quiver data, so each V_i carries a T -weight decomposition. Dinkins, Karpov and Krylov show [4, Thm. 5.6, Thm. 5.7] that when T acts on every V_i with *simple spectrum* — all weights distinct — the fixed point is encoded by a finite coloured poset, and the associated vertex function is a generating function over reverse plane partitions on that poset. They then write, in §5.4.2:

“We consider it an interesting open problem to find a combinatorial characterization of such w .”

They rule out three natural guesses. Dominant minuscule elements are too restrictive [4, Ex. 5.10]; so are minuscule elements [4, Ex. 5.11]; and so are fully commutative elements [4, Ex. 5.12]. Each example exhibits a w with simple spectrum lying outside the class.

1.2. Results. Our main theorem answers the question by constructing V explicitly, uniformly in the type, from any reduced word for w .

Let $\Pi = \Pi(\Gamma)$ be the preprojective algebra over a field, S_i the simple module at i , and I_i the injective envelope of S_i . For a Π -module M and $j \in I$ let $\text{soc}_{(j)}(M)$ be the sum of all submodules of M isomorphic to S_j , and for a sequence $\sigma = (j_1, \dots, j_p)$ let $\text{soc}_\sigma(M)$ be the iterated socle of Geiss–Leclerc–Schröer [6, §2.4], defined by $M_0 = 0$ and $M_q/M_{q-1} = \text{soc}_{(j_q)}(M/M_{q-1})$.

Definition 1.1 (the c -trace). Let $w = s_{i_1} \cdots s_{i_\ell}$ be a reduced word. Set $\mu^{(\ell)} = \lambda$ and, for $p = \ell, \ell - 1, \dots, 1$,

$$c_p = \langle \alpha_{i_p}^\vee, \mu^{(p)} \rangle, \quad \mu^{(p-1)} = \mu^{(p)} - c_p \alpha_{i_p} = s_{i_p} \mu^{(p)}.$$

The c -weighted sequence of the word is $\sigma = (i_\ell^{c_\ell}, i_{\ell-1}^{c_{\ell-1}}, \dots, i_1^{c_1})$, each letter repeated according to its multiplicity, in this order.

Note $\mu^{(0)} = w\lambda = \mu$ and $\sum_{p: i_p=i} c_p = v_i$, so the c -trace recomputes the dimension vector. The invariant $c_{\max} = \max_p c_p$ equals 1 exactly when w is λ -minuscule in the sense of [4, Def. 5.1].

Theorem 1.2 (Main Theorem). *Let Γ be simply-laced Dynkin, $\lambda = \omega_\kappa$, $w \in W$ of minimal length in wW_λ , and σ the c -weighted sequence of any reduced word for w . Then*

$$\underline{\dim} \text{soc}_\sigma(I_\kappa) = v = \lambda - w\lambda,$$

and consequently

$$V \cong \text{soc}_\sigma(I_\kappa)$$

as graded representations of the doubled quiver. In particular the isomorphism class of $\text{soc}_\sigma(I_\kappa)$ does not depend on the chosen reduced word.

The grading on the right is the path-length grading of I_κ ; Proposition 2.6 identifies it with the T -weight grading, and §5 explains how to read off the coloured poset of [4, §5.4] from it. The criterion is then immediate.

Corollary 1.3. *T acts on every V_i with simple spectrum if and only if the coloured graded set underlying $\text{soc}_\sigma(I_\kappa)$ has no repeated (colour, degree) pair, i.e. $\dim \text{soc}_\sigma(I_\kappa)_i[d] \leq 1$ for all $i \in I$ and all d .*

The construction is effective: I_κ has dimension $\lambda - w_0\lambda$, independent of w , and the socle steps are kernel computations of size bounded by the graded pieces of I_κ , so the criterion is decidable in time polynomial in $|I|$ and $\ell(w)$ from any reduced word.

Two refinements make the answer sharper. First, only one vertex needs to be inspected.

Theorem 1.4. *Let t be the trivalent node of Γ (types D and E). If $\dim V_t[d] \leq 1$ for all d , then $\dim V_i[d] \leq 1$ for all i and d .*

Theorem 1.4 follows in both types from a single inward-injectivity lemma about I_κ (Lemma 7.1, §7). In type E that lemma is a *complete* finite computation, since E_6 , E_7 and E_8 have only finitely many triples (fundamental weight, vertex, degree); in type D it is proved uniformly in the rank (Lemma 7.2). Remark 6.4 records why the more obvious argument, containment of multiplicity-2 windows, fails in both types. In types D and E this reduces the criterion to a single vertex; in type A there is no trivalent node, but there the question is settled outright:

Proposition 1.5. *If λ is minuscule then, for w of minimal length in wW_λ , $c_p = 1$ for every step of every reduced word, w is λ -minuscule, and T acts with simple spectrum. In particular the spectrum is always simple in type A , for the spin weights ω_{n-1}, ω_n of D_n , for ω_1 and ω_6 of E_6 , and for ω_7 of E_7 .*

Second, in type D the criterion becomes a statement about the signs of μ with no reference to reduced words, heaps, or fully commutative elements.

Theorem 1.6 (type D). *Let $\Gamma = D_n$, $t = n - 2$, and $\lambda = \omega_k$ with $k \leq n - 2$. Write $\mu = w\lambda = \sum_{i \in S} \eta_i e_i$ in Bourbaki coordinates, $|S| = k$, $\eta_i = \pm 1$. Put*

$$m = \#\{i \leq n - 2 : i \in S, \eta_i = -1\}, \quad r = |S \cap \{n - 1, n\}|.$$

Then T acts with simple spectrum if and only if $r + 2m \leq 2$; equivalently, iff $m = 0$, or $m = 1$ and $r = 0$. For the spin weights the spectrum is always simple. The proof is unconditional for the spin weights; for $k = 2$ it rests on a generation pattern checked for D_4 – D_7 , and for $k \geq 3$ on Proposition 6.5, which we verify computationally. See Remark 6.6.

Finally we record a counterexample in the direction opposite to [4, Ex. 5.12].

Example 1.7. In D_4 take $w = s_2 s_4 s_3 s_1 s_2$, so $v = (1, 3, 1, 1)$. This w is fully commutative, but the spectrum is *not* simple: the path-length graded dimensions of V inside I_2 are $\{2 : 1\}$ in degree 4, $\{1 : 1, 3 : 1, 4 : 1\}$ in degree 3 and $\{2 : 2\}$ in degree 2, so, read from the top down as in the diagrams of [4], the rows are $[2]$, $[1, 3, 4]$, $[2, 2]$. Together with [4, Ex. 5.12], which exhibits a w with simple spectrum that is not fully commutative, this shows that full commutativity is neither necessary nor sufficient.

1.3. Method. The proof of Theorem 1.2 is short and does not proceed case by case. It has two ingredients. The first is an exact formula for the size of a single socle step, valid for an arbitrary submodule of I_κ :

Lemma 1.8 (§3). *Let $X \subseteq I_\kappa$ be a submodule, $v' = \underline{\dim} X$, $\mu' = \lambda - \sum_i v'_i \alpha_i$. Then for every $i \in I$*

$$\dim \text{soc}_{(i)}(I_\kappa/X) = \langle \alpha_i^\vee, \mu' \rangle + \dim \text{top}_i(X),$$

where $\text{top}_i(X) = \text{Hom}_\Pi(X, S_i)$.

This comes from applying $\text{Hom}(S_i, -)$ to $0 \rightarrow X \rightarrow I_\kappa \rightarrow Q \rightarrow 0$, using $\text{Ext}^1(S_i, I_\kappa) = 0$, and then Crawley-Boevey's formula $\dim \text{Ext}_\Pi^1(M, N) = \dim \text{Hom}(M, N) + \dim \text{Hom}(N, M) - (\underline{\dim} M, \underline{\dim} N)$ [3].

The second ingredient is a squeeze. Since $\dim \text{top}_i(X) \geq 0$, each socle step grows X by at least c ; summing the c 's along a reduced word for the longest element gives exactly $\lambda - w_0 \lambda$, which is precisely $\underline{\dim} I_\kappa$. As $\text{soc}_\sigma(I_\kappa) \subseteq I_\kappa$ the two bounds meet, so *every* step contributes exactly c and every top-term vanishes. Restricting to the prefix corresponding to w gives Theorem 1.2, once uniqueness of the submodule with w -extremal dimension vector [8, Prop. 4.9] identifies the result with V .

1.4. Verification. Every statement in this paper is accompanied by machine verification against an independent solver that enumerates graded submodules directly. In particular the construction reproduces the three coloured posets printed in [4, pp. 82–83], Hasse edges included, and the criterion agrees with the solver on 2179 elements across types A , D and E — every fundamental weight of A_3 – A_6 , D_4 – D_7 , E_6 , together with $E_7 \omega_1$ and $E_7 \omega_7$, whose Weyl orbits have at most 400 elements. See §8.

1.5. Acknowledgements. The problem was communicated by V. Krylov.

2. SETUP

2.1. The preprojective algebra. Let Q be a quiver with underlying graph Γ and let $\overline{Q}$ be its double, with a^* the reverse of a . The preprojective algebra is

$$\Pi = \Pi(\Gamma) = \mathbb{k}\overline{Q} / \left(\sum_a (aa^* - a^*a) \right),$$

where $\mathbb{k}$ is any field. Every argument below is characteristic-free; Crawley-Boevey's formula, which we use in §3, holds over an arbitrary field. The computer verifications of §8 were run over $\mathbb{F}_p$ for $p = 5, 7, 11, 13$ with identical results. A Π -module is the same thing as a representation of $\overline{Q}$ satisfying the moment map equation at every vertex, which is exactly the datum carried by a point of $\mathcal{M}(v, w)$ with vanishing stability-free part. We use throughout the following standard facts, valid because Γ is Dynkin.

Proposition 2.1. $\Pi(\Gamma)$ is finite dimensional and self-injective. Its Nakayama permutation ν is determined by $\text{soc}(\Pi e_j) \cong S_{\nu(j)}$, and

$$I_\kappa \cong \Pi e_{\nu^{-1}(\kappa)}.$$

Moreover ν is the identity except in types A_n , D_{odd} and E_6 , where it is induced by the nontrivial diagram automorphism.

Proposition 2.2. $\dim I_\kappa = \omega_\kappa - w_0 \omega_\kappa$ in root coordinates.

Proposition 2.1 is classical; see for instance [3, §1]. Proposition 2.2 follows from the description of Πe_j as $\bigoplus_{d \geq 0} \tau^{-d} P_j$ over the path algebra of an orientation of Γ , whose total dimension vector is $\omega_j + \omega_{j^*} = \omega_j - w_0 \omega_j$; it is verified directly for every fundamental weight of $A_3, A_4, A_5, D_4, D_5, D_6, D_7, E_6$ and E_7 in §8.

Every Πe_j is graded by path length, and *this is the grading we use throughout*; it coincides with the socle-layer grading read in reverse. We write $M_i[d]$ for the degree- d part of M at vertex i , and $d_{\max}$ for the top degree of I_κ .

2.2. Iterated socles. For $j \in I$, $\text{soc}_{(j)}(M)$ is the largest submodule of M that is a direct sum of copies of S_j ; concretely it is the set of $x \in M_j$ annihilated by every arrow out of j . Given a sequence $\sigma = (j_1, \dots, j_p)$ define $\text{soc}_\sigma(M) = M_p$ where $M_0 = 0$ and M_q is the preimage in M of $\text{soc}_{(j_q)}(M/M_{q-1})$. We use only that soc_σ is a submodule and that $\text{soc}_\sigma(M) \subseteq \text{soc}_{\sigma'}(M)$ when σ is a prefix of σ' .

2.3. Extremal fixed points. We use Savage–Tingley's identification of the quiver variety with a quiver Grassmannian.

Proposition 2.3 ([8, Thm. 4.4, Prop. 4.9]; [1, 2]). $\mathfrak{L}(v, w) \cong \text{Gr}_v(\bigoplus_i I_i^{w_i})$, and if $v = \lambda - w\lambda$ is w -extremal then there is exactly one submodule of $\bigoplus_i I_i^{w_i}$ with dimension vector v .

The uniqueness has a sharper and more useful form. Baumann and Kamnitzer attach to a Π -module T a pseudo-Weyl polytope $\text{Pol}(T)$, and show [1] that it is the convex hull of the dimension vectors of the submodules of T ; Baumann, Kamnitzer and Tingley prove the *rigidity* property [2, §3.1]: *if x is a vertex of $\text{Pol}(T)$ then T has a unique subobject X with $\dim X = x$* . For $T = I_\kappa$ the polytope is the weight polytope of ω_κ reflected in the origin, so its vertices are exactly the w -extremal dimension vectors, and Proposition 2.3 is the vertex case. Two consequences we use below:

Corollary 2.4. Let $U \subseteq I_\kappa$ be a submodule. Then $\mu = \lambda - \dim U$ is a weight of the irreducible representation $L(\omega_\kappa)$, and

$$\delta(\dim U) = 2(\dim U)_\kappa - (\dim U)^\top C(\dim U) = (\lambda, \lambda) - (\mu, \mu) \geq 0,$$

with equality if and only if $\underline{\dim} U$ is a vertex of $\text{Pol}(I_\kappa)$, i.e. w -extremal. If I_κ carries two distinct submodules with the same dimension vector v , then v is not a vertex and $\delta(v) > 0$.

In type D_n with $\lambda = \omega_k$, $k \leq n - 2$, the weights of $L(\omega_k)$ are the μ with entries in $\{0, \pm 1\}$ and $|\text{supp } \mu| \leq k$ congruent to k modulo 2, so Corollary 2.4 gives $\delta(\underline{\dim} U) = k - |\text{supp } \mu|$ — verified on all 2014 graded submodules of I_κ for $D_5\omega_2$, $D_5\omega_3$, $D_6\omega_3$, $D_6\omega_4$, $D_7\omega_3$.

For $\lambda = \omega_\kappa$ the framing is $\bigoplus_i I_i^{w_i} = I_\kappa$, so Proposition 2.3 says: I_κ has a unique submodule of dimension vector v , and the fixed point V is that submodule (we use [8, Thm. 4.4] in the convention in which a point of the quiver Grassmannian is a submodule of the injective, not a quotient of it). Hence to identify V it suffices to exhibit any submodule of I_κ of the right dimension vector. That is what Theorem 1.2 does.

It remains to match the two gradings.

Lemma 2.5 (the bidegree is determined). Π is bigraded by the numbers of a - and a^* -arrows separately, the relations $\sum_a (aa^* - a^*a)$ being bihomogeneous of bidegree $(1, 1)$. For Γ a Dynkin diagram — hence a tree — the piece $(\Pi e_j)_i[d]$ is pure of a single bidegree, namely

$$p = \frac{1}{2}(d + f(i) - f(j)), \quad q = \frac{1}{2}(d - f(i) + f(j)),$$

where f is any function on the vertices with $f(\text{head}) - f(\text{tail}) = 1$ on every arrow of the chosen orientation.

Proof. Such an f exists because Γ is a tree: define it by walking out from any root; there is no cycle to obstruct consistency. An a -arrow raises f by 1 and its reverse lowers it by 1, so along any path from j to i one has $p - q = f(i) - f(j)$. With $p + q = d$ this determines p and q . $\square$

Proposition 2.6. The T -weight decomposition of V_i coincides with its path-length decomposition, up to an overall affine normalisation independent of i . In particular

$$T \text{ acts on every } V_i \text{ with simple spectrum} \iff \dim V_i[d] \leq 1 \text{ for all } i \text{ and } d.$$

Proof. Because $\lambda = \omega_\kappa$ is fundamental the framing space is one dimensional, so the framing torus acts on all of V through a single character and contributes only a global shift. The rest of T acts through the bigrading of Π : any such torus assigns to $(\Pi e_{\kappa'})_i[d]$ a weight that is a fixed integral combination $\alpha p + \beta q$ of the two bidegrees. By Lemma 2.5 both p and q are affine functions of d with i, κ' fixed, with p and q each of slope $\pm \frac{1}{2}$ in d ; hence $\alpha p + \beta q$ is an affine function of d of slope $\frac{1}{2}(\alpha + \beta)$, which is injective in d unless $\alpha + \beta = 0$. The case $\alpha + \beta = 0$ is excluded because the symplectic form has nonzero weight, that being what makes the fixed locus of T on $\mathcal{M}(v, w)$ isolated. So the weight determines and is determined by d , degree by degree.

Finally V is a graded submodule: since v is w -extremal, Proposition 2.3 gives a unique submodule of I_κ with dimension vector v , and uniqueness forces it to be stable under the whole torus, hence to carry the restricted grading. $\square$

3. THE SOCLE-STEP IDENTITY

Proof of Lemma 1.8. Write $I = I_\kappa$ and let $Q = I/X$, so that $0 \rightarrow X \rightarrow I \rightarrow Q \rightarrow 0$ is exact. Applying $\text{Hom}_\Pi(S_i, -)$ gives

$$0 \rightarrow \text{Hom}(S_i, X) \rightarrow \text{Hom}(S_i, I) \rightarrow \text{Hom}(S_i, Q) \rightarrow \text{Ext}^1(S_i, X) \rightarrow \text{Ext}^1(S_i, I).$$

The last term vanishes because I is injective. Since $\text{soc}(I_\kappa) = S_\kappa$ we have $\dim \text{Hom}(S_i, I) = \delta_{i\kappa}$, and $\dim \text{Hom}(S_i, Q) = \dim \text{soc}_{(i)}(Q)$. Therefore

$$\dim \text{soc}_{(i)}(Q) = \delta_{i\kappa} - \dim \text{soc}_{(i)}(X) + \dim \text{Ext}^1(S_i, X). \quad (1)$$

By Crawley-Boevey's formula for modules over a preprojective algebra [3, Lem. 1],

$$\dim \operatorname{Ext}_{\Pi}^1(M, N) = \dim \operatorname{Hom}(M, N) + \dim \operatorname{Hom}(N, M) - (\underline{\dim} M, \underline{\dim} N),$$

where $(\cdot, \cdot)$ is the symmetrised Tits form, so that $(\underline{\dim} M, \underline{\dim} N) = (\underline{\dim} M)^\top C (\underline{\dim} N)$. Taking $M = S_i$ and $N = X$, and noting $\operatorname{Hom}(S_i, X) = \operatorname{soc}_{(i)}(X)$ and $\operatorname{Hom}(X, S_i) = \operatorname{top}_i(X)$,

$$\dim \operatorname{Ext}_{\Pi}^1(S_i, X) = \dim \operatorname{soc}_{(i)}(X) + \dim \operatorname{top}_i(X) - (\alpha_i, v').$$

Substituting into (1) the socle terms cancel and

$$\dim \operatorname{soc}_{(i)}(Q) = \delta_{i\kappa} - (\alpha_i, v') + \dim \operatorname{top}_i(X).$$

Finally $(\alpha_i, v') = \sum_j a_{ij} v'_j = \langle \alpha_i^\vee, \sum_j v'_j \alpha_j \rangle$ and $\langle \alpha_i^\vee, \omega_\kappa \rangle = \delta_{i\kappa}$, so $\delta_{i\kappa} - (\alpha_i, v') = \langle \alpha_i^\vee, \mu' \rangle$. $\square$

Corollary 3.1. *A socle step of colour i applied to $X \subseteq I_\kappa$ enlarges X by exactly $\langle \alpha_i^\vee, \mu' \rangle + \dim \operatorname{top}_i(X)$ dimensions, all at vertex i . In particular the growth is at least $\langle \alpha_i^\vee, \mu' \rangle$, with equality if and only if S_i does not occur in the head of X .*

Remark 3.2. Lemma 1.8 also explains why repeating a colour is harmless. Suppose a step of colour i with $c = \langle \alpha_i^\vee, \mu' \rangle$ has just added exactly c dimensions, producing X' . Those c new elements lie in the head of X' , so $\dim \operatorname{top}_i(X') = c$; and the weight has become $s_i \mu'$, with $\langle \alpha_i^\vee, s_i \mu' \rangle = -c$. Lemma 1.8 gives $\dim \operatorname{soc}_{(i)}(I_\kappa/X') = -c + c = 0$: further repeats add nothing. Thus $\operatorname{soc}_\sigma$ is unchanged if each block i^c is replaced by $i^{c'}$ for any $c' \geq c$, and the c -weighted sequence may equivalently be described as “apply $\operatorname{soc}_{(i)}$ until it stabilises”.

4. PROOF OF THE MAIN THEOREM

Lemma 4.1. *Let w be of minimal length in wW_λ and $\mu = w\lambda$.*

- (i) *Every c -trace value of every reduced word for w satisfies $c_p \geq 1$.*
- (ii) *If $\mu \neq w_0\lambda$ there is a simple reflection s_i with $\langle \alpha_i^\vee, \mu \rangle \geq 1$, and for any such i one has $\ell(s_i w) = \ell(w) + 1$ with $s_i w$ again of minimal length in $s_i wW_\lambda$.*

Consequently every reduced word for w extends on the left to a reduced word for the longest minimal-length representative w_0^λ , with all c -trace values ≥ 1 .

Proof. (i) Write $u = s_{i_{p+1}} \cdots s_{i_\ell}$, so $\mu^{(p)} = u\lambda$. The word is reduced, so $\ell(s_{i_p} u) = \ell(u) + 1$ and hence $u^{-1} \alpha_{i_p} > 0$; since λ is dominant, $c_p = \langle (u^{-1} \alpha_{i_p})^\vee, \lambda \rangle \geq 0$. It remains to exclude $c_p = 0$. If $c_p = 0$ then $s_{i_p} u\lambda = u\lambda$, hence $w\lambda = s_{i_1} \cdots s_{i_{p-1}} u\lambda$, and $s_{i_1} \cdots s_{i_{p-1}} u$ has length at most $\ell - 1$. This contradicts the minimality of $\ell(w)$ among elements sending λ to μ .

(ii) If $\langle \alpha_i^\vee, \mu \rangle \leq 0$ for every i then μ is antidominant, hence $\mu = w_0\lambda$. So assume $\langle \alpha_i^\vee, w\lambda \rangle > 0$. Then $\langle \lambda, w^{-1} \alpha_i^\vee \rangle > 0$; since λ is dominant this forces $w^{-1} \alpha_i > 0$, and therefore $\ell(s_i w) = \ell(w) + 1$. Let $J = \{j : \langle \alpha_j^\vee, \lambda \rangle = 0\}$, so that $W_\lambda = W_J$ and w is minimal in wW_J if and only if $w\alpha_j > 0$ for all $j \in J$. For $j \in J$ we have $w\alpha_j > 0$, and $s_i(w\alpha_j) < 0$ could only happen if $w\alpha_j = \alpha_i$; but then $\langle \alpha_i^\vee, w\lambda \rangle = \langle \alpha_j^\vee, \lambda \rangle = 0$, contradicting $\langle \alpha_i^\vee, w\lambda \rangle > 0$. Hence $s_i w\alpha_j > 0$ for all $j \in J$ and $s_i w$ is minimal in its coset. $\square$

Proof of Theorem 1.2. Fix a reduced word for w and extend it, by Lemma 4.1, to a reduced word for w_0^λ ; let σ_{full} be its c -weighted sequence and σ the prefix belonging to w . Let $X_0 = 0 \subseteq X_1 \subseteq \cdots \subseteq X_L = \operatorname{soc}_{\sigma_{\text{full}}}(I_\kappa)$ be the resulting filtration, indexed by the blocks $i_p^{c_p}$ of σ_{full} rather than by individual letters.

By Corollary 3.1 applied to the first letter of the p -th block, and because the remaining letters of the block can only enlarge X further,

$$\dim X_p - \dim X_{p-1} \geq c_p \quad (1 \leq p \leq L). \quad (2)$$

Here $c_p \geq 1$ by Lemma 4.1(i); it is this positivity that makes (2) a lower bound strong enough to be squeezed. Summing, and using that the c -trace telescopes, $\sum_p c_p \alpha_{i_p} = \lambda - w_0^\lambda \lambda = \lambda - w_0 \lambda$, we get

$$\dim \text{soc}_{\sigma_{\text{full}}}(I_\kappa) \geq \sum_p c_p = \dim(\lambda - w_0 \lambda) = \dim I_\kappa,$$

the last equality by Proposition 2.2. On the other hand $\text{soc}_{\sigma_{\text{full}}}(I_\kappa) \subseteq I_\kappa$, so the inequality is an equality. Consequently every inequality in (2) is an equality, and since the p -th block only enlarges X at vertex i_p ,

$$\underline{\dim} X_p - \underline{\dim} X_{p-1} = c_p \alpha_{i_p} \quad \text{for all } p.$$

Truncating at the prefix σ gives $\underline{\dim} \text{soc}_\sigma(I_\kappa) = \sum_{p \leq \ell} c_p \alpha_{i_p} = \lambda - w\lambda = v$. By Proposition 2.3 the submodule of I_κ with this dimension vector is unique and equals V . Word-independence follows, since the right-hand side depends only on w . $\square$

Corollary 4.2. *At the first letter of each block i^c of σ one has $\text{top}_i(X) = 0$ and $\dim \text{soc}_{(i)}(I_\kappa/X) = c$ exactly. (Within a block this fails, and must: after the first letter $\text{top}_i = c$ and $\text{soc}_{(i)}(I_\kappa/X) = 0$, as computed in Remark 3.2.)*

Remark 4.3. Corollary 4.2 has a purely combinatorial shadow. Only reflections at i and at its neighbours change $\langle \alpha_i^\vee, \cdot \rangle$, and a colour- i step flips that pairing from c to $-c$; so between two consecutive colour- i blocks of a c -trace some neighbouring colour must intervene. It is exactly that intervening step which absorbs S_i out of the head of X .

Proof of Proposition 1.5. If λ is minuscule then $\langle \alpha^\vee, \nu \rangle \in \{0, \pm 1\}$ for every root α and every weight ν of $L(\lambda)$, so every c_p equals 1 and w is λ -minuscule in the sense of [4, Def. 5.1]. By Stembridge [9, Prop. 2.1, §2] a λ -minuscule element is fully commutative, so its heap $H(w)$ is independent of the reduced word; and λ -minuscule heaps are ranked, with each colour fibre $H(w)_i$ a chain. By [5, Prop. 3.8, Lem. 3.10] the fixed point representation is the heap module $\mathbb{C}H(w)$, with V_i spanned by $H(w)_i$. The rank function of $H(w)$ is the path-length degree of Proposition 2.6, and a chain occupies pairwise distinct ranks, so every graded piece of V_i has dimension at most one. $\square$

5. READING THE COLOURED POSET

Let $X = \text{soc}_\sigma(I_\kappa) \cong V$ and assume the criterion of Corollary 1.3 holds, so that $\dim X_i[d] \leq 1$ for all i, d ; only then are the vertices of the poset canonically labelled. Define a coloured graded set

$$P(w) = \{(i, d, \text{basis vector of } X_i[d])\},$$

coloured by i and ranked by $d_{\max} - d$, with $(i, d) < (j, d+1)$ whenever j is adjacent to i and the corresponding arrow of Π is nonzero on the relevant vectors. Reading the degrees in decreasing order recovers the Hasse diagrams of [4, §5.4], whose top row is the top degree $d_{\max}$. Then [4, Thm. 5.6, Thm. 5.7] express the vertex function as a sum over reverse plane partitions on $P(w)$.

Example 5.1 ([4, Ex. 5.10]). $\Gamma = D_4$, $\lambda = \omega_2$, $w = s_1 s_2 s_3 s_4 s_2$. The c -trace gives $v = (2, 2, 1, 1)$ and $\sigma = (2, 4, 3, 2, 1, 1)$. The construction returns rows $[2], [1, 3, 4], [2], [1]$ read from the top degree down, matching the diagram on p. 82 of [4] including its six covering relations.

Example 5.2 ([4, Ex. 5.11, Ex. 5.12]). $\Gamma = D_5$, $\lambda = \omega_2$. For $w = s_2 s_1 s_3 s_4 s_5 s_3 s_2$ the construction returns $[2], [1, 3], [2, 4, 5], [3], [2]$, and for $w = s_2 s_1 s_2 s_3 s_4 s_5 s_3 s_2$ it returns $[2], [1, 3], [2, 4, 5], [3], [2], [1]$, both matching p. 83 of [4] with all covering relations.

6. TYPE D

Throughout this section $\Gamma = D_n$ with Bourbaki labelling $\alpha_i = e_i - e_{i+1}$ for $i \leq n-1$ and $\alpha_n = e_{n-1} + e_n$, so that the trivalent node is $t = n-2$.

Lemma 6.1. *Let $\lambda = \omega_k$ with $k \leq n-2$ and $\mu = w\lambda = \sum_{i \in S} \eta_i e_i$. With m and r as in Theorem 1.6 one has $v_t = r + 2m$.*

Proof. Write $\lambda - \mu = \sum_j c_j e_j$. Comparing coefficients of e_j in $\lambda - \mu = \sum_i v_i \alpha_i$ gives $c_1 = v_1$ and $c_j = v_j - v_{j-1}$ for $2 \leq j \leq n-2$, so $v_t = \sum_{j \leq n-2} c_j$. Since $k \leq n-2$ the support of λ lies in the first $n-2$ coordinates, whence $\sum_{j \leq n-2} c_j = k - \sum_{j \leq n-2} \mu_j$. Writing p and m for the number of positive and negative signs of μ among the first $n-2$ coordinates, $\sum_{j \leq n-2} \mu_j = p - m$ and $p + m = k - r$, so $v_t = k - (p - m) = r + 2m$. $\square$

Lemma 6.2 (the recursion). *Let $a_d(i) = \dim(I_\kappa)_i[d]$. Then for every i and every d strictly below the top degree,*

$$a_{d+1}(i) = \sum_{j \sim i} a_d(j) - a_{d-1}(i), \quad a_0(i) = \delta_{i, \kappa'}, \quad a_{-1} = 0,$$

where $\kappa' = \nu^{-1}(\kappa)$.

Proof. By construction the degree- $(d+1)$ part of the free module at j is $\bigoplus_{i \sim j} (I_\kappa)_i[d]$, and the preprojective relations impose one condition for each basis vector of $(I_\kappa)_j[d-1]$. These conditions are independent precisely when $Y : (I_\kappa)_j[d-1] \rightarrow \bigoplus_{i \sim j} (I_\kappa)_i[d]$ is injective. A nonzero element of $\ker Y$ would span a copy of S_j inside I_κ ; but I_κ has simple socle S_κ , concentrated in the top degree. Hence Y is injective below the top degree. $\square$

Lemma 6.3 (method of images). *For $\Gamma = D_n$ and $\lambda = \omega_k$ with $k \leq n-2$, the colour-t multiplicity sequence of I_κ is $(1, 2, \dots, 2, 1)$ with $k-1$ twos, occupying the $k+1$ degrees $t-k, t-k+2, \dots, t+k$. More generally vertex i carries multiplicity 2 in exactly $\max(0, i+k-n+1)$ degrees and multiplicity ≤ 2 throughout.*

Proof. The datum $a_0 = \delta_k$ is fixed by the diagram automorphism exchanging $n-1$ and n , hence so is the solution: $a_d(n-1) = a_d(n) =: b_d$. Setting $a_d(t+1) := 2b_d$, the recursion of Lemma 6.2 becomes the free chain recursion $a_{d+1}(i) + a_{d-1}(i) = a_d(i-1) + a_d(i+1)$ on $1, \dots, L$ with $L = t+1 = n-1$, subject to a Dirichlet condition $a_d(0) = 0$ and a Neumann (mirror) condition $a_d(L+1) = a_d(L-1)$.

On the infinite chain the solution with $u_0 = \delta_k, u_{-1} = 0$ is $u_d(i) = 1$ if $|i-k| \leq d$ and $i \equiv k+d \pmod{2}$, and 0 otherwise (induction on d). The Dirichlet condition is an odd reflection about 0 and the mirror condition an even reflection about L ; the group they generate is infinite dihedral, with images of k at $2Lj \pm k$ carrying alternating signs. Hence

$$a_d(i) = \sum_{x \text{ image of } k} \text{sign}(x) \cdot [|i-x| \leq d \text{ and } i \equiv x+d \pmod{2}].$$

The four images nearest t are k and $2L-k$ with sign $+$, at distances $t-k$ and $t+2-k$, and $-k$ and $2L+k$ with sign $-$, at distances $t+k$ and $t+2+k$. (The image solution vanishes identically once d exceeds $t+2+k$, and the total $\sum_{d,i} a_d(i)$ it predicts equals $\lambda - w_0 \lambda = \underline{\dim} I_\kappa$ by Proposition 2.2; this pins the top degree and shows that the recursion of Lemma 6.2, valid below the top degree, determines I_κ completely.) They enter the window one at a time as d grows, so $a_d(t)$ is 0 for $d < t-k$, then 1, then 2 on $[t+2-k, t+k)$, then 1 on $[t+k, t+2+k)$, then 0; the degrees of the correct parity in the middle range number $(2k-4)/2 + 1 = k-1$. The same computation at a vertex $i \leq t$ of the chain shows the value 2 occurs exactly for

$\max(|i - k|, 2L - k - i) \leq d < \min(i + k, 2L + k - i)$, giving $\max(0, i + k - n + 1)$ degrees. (At the two leaves $n - 1, n$ the folding gives $a_d(n - 1) = a_d(n) = b_d$ with $2b_d = a_d(t + 1) \leq 2$, so $b_d \leq 1$ and they never carry multiplicity 2; the displayed count applies only on the chain $i \leq t$.) $\square$

Remark 6.4 (window containment fails, and why). It is tempting to argue Theorem 1.4 in type D from Lemma 6.3 alone: the multiplicity-2 window of vertex $i \leq t$ is $[2L - k - i, i + k]$ with $L = n - 1$, whose endpoints satisfy $2L - k - i \geq L - k + 1$ and $i + k \leq (L - 1) + k$, so it is contained in the window of t *as an interval*. This is not enough, for two independent reasons.

First, the degrees carried by vertex i obey the parity constraint $d \equiv i + k \pmod{2}$, so the actual degree *sets* need not be nested when $i \not\equiv t \pmod{2}$. In D_6 with $\kappa = 3$ and $t = 4$, vertex 3 carries multiplicity 2 exactly at degree 4 while vertex t carries it exactly at degrees 3, 5; the sets are disjoint. (The same happens in D_7 , $k = 3$.) Second, even granting containment, knowing that a repeat at i occurs at a degree in the t -window does not produce a repeat at t : nothing transports $\dim V_i[d] \geq 2$ to $\dim V_t[d'] \geq 2$.

The mechanism that does work is inward injectivity, Lemma 7.1 below, which moves a repeat one vertex towards t and shifts the degree by one — and with that shift the D_6 numbers do line up, $i = 3$ at $d = 4$ mapping to t at $d = 5$. Lemma 7.1 is stated and used in §7; it is what proves Theorem 1.4 in both types.

Proposition 6.5 (top truncation). *Let $N \subseteq I_\kappa$ be a graded submodule with $\dim N = v = \lambda - w\lambda$ for some $w \in W$. Then the colour- t multiplicities of N are the greedy top truncation of the sequence of Lemma 6.3 to total v_t : the same degrees, filled from the highest downward.*

Note that the sequence $(1, 2^{k-1}, 1)$ of Lemma 6.3 is palindromic, so its top truncations and its bottom truncations have the same multiset of multiplicities, namely

$$v_t = 0 \mapsto (), \quad 1 \mapsto (1), \quad 2 \mapsto (1, 1), \quad 3 \mapsto (1, 2), \quad 4 \mapsto (1, 2, 1), \quad 5 \mapsto (1, 2, 2), \dots$$

Only the multiset is used below.

Proof for $k = 2$, checked for $n \leq 7$. Here the colour- t profile of I_κ is $(1, 2, 1)$ at degrees $d_0, d_0 + 2, d_0 + 4$ with $d_0 = n - 4$. Computing the submodule generated by a single colour- t element shows that its colour- t content is all of $(1, 2, 1)$ if the element sits in degree d_0 , is $(1, 1)$ in degrees $d_0 + 2, d_0 + 4$ if it sits in degree $d_0 + 2$ (for every choice of vector there), and is (1) if it sits in degree $d_0 + 4$. Hence the only colour- t profiles a submodule can have are $(0, 0, 0)$, $(0, 0, 1)$, $(0, 1, 1)$, $(0, 2, 1)$, $(1, 2, 1)$, and each is a top truncation. No hypothesis on v is needed. The generation pattern was computed for D_4 – D_7 ; it is uniform in n , but we have not written the induction out, so Theorem 1.6 at $k = 2$ carries that caveat. $\square$

Remark 6.6 (what Proposition 6.5 rests on). Call a graded submodule $M \subseteq I_\kappa$ a *staircase* if its colour- t multiplicities are the greedy top truncation of those of I_κ — zero below one degree, proper at it, full above — and if at each of the two leaves M occupies a final segment of degrees. Everything left open in this paper is the following single statement.

Conjecture 6.7. *If $\dim M$ is w -extremal then M is a staircase. Equivalently, by the rigidity property of [2] quoted in §2, a submodule that is not a staircase satisfies $\delta(\dim M) > 0$, and so is never the unique submodule of its dimension vector.*

Conjecture 6.7 gives Proposition 6.5 at once: an extremal v carries a unique submodule, namely $V = \text{soc}_\sigma(I_\kappa)$ by Theorem 1.2, and the conjecture makes it a staircase. It is verified on all 1124 graded submodules of I_κ over D_4 – D_7 : exactly 78 fail to be staircases and every one has $\delta = 2$ (`master.py`).

Proposition 6.8. *Conjecture 6.7 holds whenever $v_{n-1} = v_n$.*

Proof. Suppose M is not a staircase at t , say $M_t[d] \neq 0$ while $M_t[d+2] \subsetneq (I_\kappa)_t[d+2]$. Every nonzero $x \in M_t[d]$ is then *special* — annihilated by an arrow from t to a leaf — since otherwise the three double paths $a_{j \rightarrow t} a_{t \rightarrow j} x$, which satisfy one linear relation and hence span at most two dimensions, would by direct computation span all of $(I_\kappa)_t[d+2]$, forcing it into M . The special vectors form two lines rather than a subspace, so $M_t[d]$ is one of them, and $\dim(I_\kappa)_t[d] = 2$ at every such degree. The diagram automorphism τ exchanging the leaves fixes κ (as $k \leq n-2$), hence acts on I_κ , and exchanges the two arrows out of t to the leaves, hence the two special lines; so $\tau(M) \neq M$. When $v_{n-1} = v_n$ we have $\underline{\dim} \tau(M) = \underline{\dim} M$, and two distinct submodules of one dimension vector contradict rigidity, so $\underline{\dim} M$ is not extremal and $\delta > 0$ by Corollary 2.4. $\square$

Three further reductions are unconditional and may help. Running the induction along the socle steps, with D the colour- t threshold of the partial module X : (a) a colour- t step contributes nothing in degrees $< D-2$; (b) if it contributes at $D-2$ then $\dim X_t[D] = 1$ and the $D-2$ contribution is exactly one special line, so a single further dimension at D suffices. Both follow from the simple socle of I_κ together with the propagation of Lemma 7.2. And (c), the condition “every neighbour of t is full at degree $D+1$ ” — which is what makes a colour- t socle step reach the threshold D — is provable at *all three* neighbours, assuming only that X is an up-set at those neighbours. At a colour- t step $c = \mu'_t - \mu'_{n-1} \geq 1$ with $\mu'_t \leq 1$ forces $\mu'_{n-1} \leq 0$, so $r' = |\text{supp } \mu' \cap \{n-1, n\}|$ is 0 or 2; and if $X_t[D] \neq 0$ then $v'_t = 2a$ is even with $D = t + k - 2a$.

At the leaves, the ε -coordinate identities give

$$v'_{n-1} + v'_n = v'_t - \mu'_{n-1} \in \{2a, 2a+1\}, \quad |v'_{n-1} - v'_n| = |\mu'_n| \leq 1,$$

so $\min(v'_{n-1}, v'_n) \geq a$, which is exactly the condition for an up-set at a leaf to be full at $D+1$. At the chain neighbour, $c \geq 1$ gives $v'_{t-1} \geq 2a - [k=t]$, while the profile of I_κ satisfies

$$\sum_{d \geq D+1} \dim(I_\kappa)_{t-1}[d] \leq 2a - [k=t]$$

(checked for all D_5 – D_8 and all $k \leq n-2$, 52/52; it is a statement about the image formula of Lemma 6.3 alone), so an up-set at $t-1$ is full at $D+1$ as well. Both halves are confirmed directly, at 1104 leaf instances and 527 chain instances.

So the conjecture reduces to the up-set property at t and at its neighbours holding *simultaneously* along the construction — which is why we state it as one statement about staircases rather than as a chain of separate claims.

Two facts delimit the search. Two facts delimit the search. The conjecture genuinely needs extremality: the staircase property fails for 78 of the 1124 graded submodules above, and the analogous statement for arbitrary submodules is false. And the obvious mechanism is not available: when $X_t[D]$ is a line it is *always* one of the two special lines (527/527), so one leaf arrow annihilates it and that leaf’s fullness at $D+1$ comes from elsewhere in X .

Proof of Theorem 1.6. For the spin weights λ is minuscule and Proposition 1.5 applies. Let $k \leq n-2$. By Lemma 6.3 the colour- t profile of I_κ is $(1, 2^{k-1}, 1)$. By Theorem 1.2, V is a submodule of I_κ with $\underline{\dim} V = v$, so Proposition 6.5 applies to it and the colour- t multiplicities of V are the truncation displayed above. Every multiplicity is ≤ 1 exactly when $v_t \leq 2$. By Theorem 1.4, simplicity at t implies simplicity everywhere. Lemma 6.1 identifies $v_t = r + 2m$. $\square$

Remark 6.9. The disjunctive form “ $c_{\max} = 1$ or $v_t \leq 2$ ” of the criterion, natural from the c -trace, collapses once the spin weights are separated: for non-spin λ in type D , $c_{\max} = 1$ forces $(r, m) \in \{(0, 0), (0, 1), (1, 0), (2, 0)\}$, all of which satisfy $r + 2m \leq 2$.

7. TYPE E

Theorem 1.2 and Corollary 1.3 are uniform and therefore settle type E as well: the criterion is decidable from any reduced word. Two further points deserve comment.

7.1. Reduction to the trivalent node. As Remark 6.4 explains, no containment-of-windows argument proves Theorem 1.4, in either type: in E_6 with $\kappa = 3$ (so $I_3 = \Pi e_5$) vertex 3 carries multiplicity 2 in path-length degree 4 while vertex 4 is zero there, and the same phenomenon already occurs in D_6 with $\kappa = 3$. The mechanism that does work is the following.

For $i \neq t$ let $\text{nx}(i)$ denote the neighbour of i one step closer to t .

Lemma 7.1 (inward injectivity). *Let Γ be E_6 , E_7 or E_8 . For every fundamental weight κ , every $i \neq t$ and every degree d with $\dim(I_\kappa)_i[d] \geq 2$, the arrow*

$$(I_\kappa)_i[d] \longrightarrow (I_\kappa)_{\text{nx}(i)}[d+1]$$

is injective.

Proof. The three diagrams have finitely many triples (κ, i, d) in total, so this is a finite verification: 11 instances in E_6 , 55 in E_7 and 249 in E_8 , all injective. See `inwardlemma.py`. $\square$

Proof of Theorem 1.4 in type E . Suppose $\dim V_i[d] \geq 2$ for some $i \neq t$. Then $\dim(I_\kappa)_i[d] \geq \dim V_i[d] \geq 2$, so by Lemma 7.1 the inward arrow is injective on $(I_\kappa)_i[d]$, hence on the subspace $V_i[d]$. Since V is a submodule of I_κ , the image lies in $V_{\text{nx}(i)}[d+1]$, so $\dim V_{\text{nx}(i)}[d+1] \geq 2$. Iterating $\text{dist}(i, t)$ times gives $\dim V_t[d + \text{dist}(i, t)] \geq 2$. Contrapositively, if V_t is multiplicity-free then so is V . $\square$

In type D the same statement holds, and there it can be proved outright, uniformly in the rank.

Lemma 7.2 (Lemma 7.1 in type D). *Let $\Gamma = D_n$, $t = n - 2$, let $i < t$ be a chain vertex and let d satisfy $\dim(I_\kappa)_i[d] \geq 2$. Then the inward arrow $(I_\kappa)_i[d] \rightarrow (I_\kappa)_{i+1}[d+1]$ is injective.*

Proof. Suppose $x \neq 0$ lies in the kernel. For $i < t$ the vertex i has exactly the two neighbours $i-1$ and $i+1$. By Lemma 6.2 the socle of I_κ is simple and concentrated in the top degree, so an element killed by every arrow out of i is either zero or that socle generator; since $\dim(I_\kappa)_i[d] \geq 2$ we are below the top degree, and therefore

$$y := a_{i \rightarrow i-1}(x) \neq 0.$$

The preprojective relation at i reads $\varepsilon a_{i+1 \rightarrow i} a_{i \rightarrow i+1} + \varepsilon' a_{i-1 \rightarrow i} a_{i \rightarrow i-1} = 0$ on $(I_\kappa)_i[d]$; applied to x , whose first term vanishes by hypothesis, it gives $a_{i-1 \rightarrow i}(y) = 0$. So y is a nonzero element of $(I_\kappa)_{i-1}[d+1]$ killed by its inward arrow: the same situation, one vertex further from t and one degree higher.

Iterating $i-1$ times produces a nonzero $z \in (I_\kappa)_1[d+i-1]$ killed by the arrow $1 \rightarrow 2$. In D_n the vertex 1 has 2 as its only neighbour, so z is killed by every arrow out of 1 and hence lies in $\text{soc}(I_\kappa) = S_\kappa$, forcing $\kappa = 1$. But by Lemma 6.3 a chain vertex i carries multiplicity 2 in exactly $\max(0, i+k-n+1)$ degrees, which vanishes for $k=1$ and $i \leq n-2$; so $\dim(I_\kappa)_i[d] \geq 2$ is impossible when $\kappa = 1$. Contradiction. $\square$

Remark 7.3. The two leaves $n-1, n$ need no argument: by the folding in Lemma 6.3 they satisfy $a_d(n-1) = a_d(n) = b_d$ with $2b_d = a_d(t+1) \leq 2$, so $b_d \leq 1$ and Lemma 7.2 is vacuous there. With Lemma 7.2 in hand the proof of Theorem 1.4 given above applies verbatim in type D , for every rank. It is also confirmed directly for D_4 – D_{10} (126 instances) by `inwardlemma.py`, and the two inputs — simple socle and the relation at a chain vertex — by `inwardproof.py`.

7.2. No closed form of type- D shape. What does *not* carry over is the closed form of Theorem 1.6: v_t is no longer a complete invariant. In E_6 with $\lambda = \omega_3$ both simple and non-simple elements occur at $v_t = 3$; likewise in E_7 with $\lambda = \omega_2$ at $v_t = 4$. A closed form must therefore involve more than the multiplicity at the trivalent node. We have not found one.

The colour- t multiplicity profiles of I_κ in type E are finite data, computed once and for all; in E_6 they are $(1, 1, 1, 1)$, $(1, 1, 2, 1, 1)$, $(1, 2, 2, 2, 1)$, $(1, 2, 3, 3, 2, 1)$, $(1, 2, 2, 2, 1)$, $(1, 1, 1, 1)$ for $k = 1, \dots, 6$, and the maximum multiplicity anywhere in I_κ is 6, attained in E_8 at $k = 4$. This is also why the truncation mechanism of §6 is specific to type D : it turns on $\dim(I_\kappa)_t[d] \leq 2$, so that the three double paths at t , subject to one relation, can span everything.

8. VERIFICATION

The accompanying code reproduces every claim; each script runs standalone and prints a pass/fail line.

- `proof.py` — the five steps of §4. The Crawley-Boevey identity holds at 5517/5517 socle steps; $\dim I_\kappa = \lambda - w_0\lambda$ at 47/47 injectives across $A_3, A_4, A_5, D_4, D_5, D_6, D_7, E_6, E_7$; $\dim \text{soc}_\sigma(I_\kappa) = v$ for 706/706 elements; the minuscule statement for 256/256; Proposition 1.4 for 981/981 elements, of which 681 are of type E and the rest of type D .
- `socdim.py` — Theorem 1.2 on 1149 elements and 12829 individual socle steps, no exceptions.
- `sweepall.py` — the aggregate criterion-versus-solver sweep quoted above; DKKCAP raises the orbit-size cutoff and DKKP changes the characteristic.
- `verify.py` — the printed posets of [4, pp. 82–83] reproduced with their Hasse edges; the criterion against an independent graded-submodule solver; word-independence; Theorem 1.6.
- `truncetest.py` — Proposition 6.5 (top truncation) on 1000 elements of D_4 – D_7 , every non-spin fundamental weight ($k \leq n-2$) with orbit at most 400, no exceptions.
- `inwardlemma.py` — Lemma 7.1, exhaustively in E_6, E_7, E_8 (315 instances) and for D_4 – D_{10} (126 instances).
- `inwardproof.py` — the inputs to Lemma 7.2: simple socle in the top degree (39 injectives, D_4 – D_9), multiplicity-freeness on the chain for $\kappa = 1$, and the relation at a chain vertex (550 instances).
- `bigrade.py` — Lemma 2.5: Γ is a tree and the potential f exists, for A_2 – A_8, D_4 – D_9, E_6, E_7, E_8 .
- `master.py` — Conjecture 6.7 over all 1124 graded submodules of I_κ for D_4 – D_7 : exactly 78 non-staircases, every one with $\delta = 2$.
- `filltop.py`, `claimG.py`, `epscrip.py`, `violators.py` — the unconditional reductions (a), (b), (c) of Remark 6.6 and the structure of the non-staircases.
- `squeeze.py`, `homological.py`, `headrule.py`, `lemA.py`, `eproject.py` — the individual ingredients.

Aggregate agreement of Corollary 1.3 with the independent solver is 2179/2179, broken down as type A 214, type D 1232, type E 733; this is the output of `sweepall.py`, which enumerates every fundamental weight of A_3 – A_6, D_4 – $D_7, E_6, E_7\omega_1$ and $E_7\omega_7$ whose Weyl orbit

has at most 400 elements and compares the criterion against the solver element by element. All computations are over a finite prime field; the whole suite was rerun at $p = 5, 11, 13$ with identical results.

REFERENCES

- [1] P. Baumann and J. Kamnitzer, *Preprojective algebras and MV polytopes*, Represent. Theory **16** (2012), 152–188.
- [2] P. Baumann, J. Kamnitzer and P. Tingley, *Affine Mirković–Vilonen polytopes*, Publ. Math. IHÉS **120** (2014), 113–205.
- [3] W. Crawley-Boevey, *On the exceptional fibres of Kleinian singularities*, Amer. J. Math. **122** (2000), 1027–1037.
- [4] A. Dinkins, D. Karpov and V. Krylov, *The quantum Hikita conjecture via quasimaps*, arXiv:2608.16746.
- [5] A. Dranowski et al., *Heaps, crystals, and preprojective algebra modules*, arXiv:2202.02490.
- [6] C. Geiss, B. Leclerc and J. Schröer, *Kac-Moody groups and cluster algebras*, Adv. Math. **228** (2011), 329–433.
- [7] H. Nakajima, *Quiver varieties and Kac-Moody algebras*, Duke Math. J. **91** (1998), 515–560.
- [8] A. Savage and P. Tingley, *Quiver Grassmannians, quiver varieties and the preprojective algebra*, Pacific J. Math. **251** (2011), 393–429.
- [9] J. Stembridge, *Minuscule elements of Weyl groups*, J. Algebra **235** (2001), 722–743.